

Deep Learning for Chest X-Ray Classification

A Technical Implementation with PyTorch Lightning


Project Overview

- Multi-label classification of chest X-rays
- 14 different pathological conditions
- Custom CNN and transfer learning approaches
- MLFlow for full experiment tracking and reproducibility
- Explainable AI techniques (Grad-CAM, Integrated Gradients)
- ❌ Didn't have time to implement hyperparameter optimization

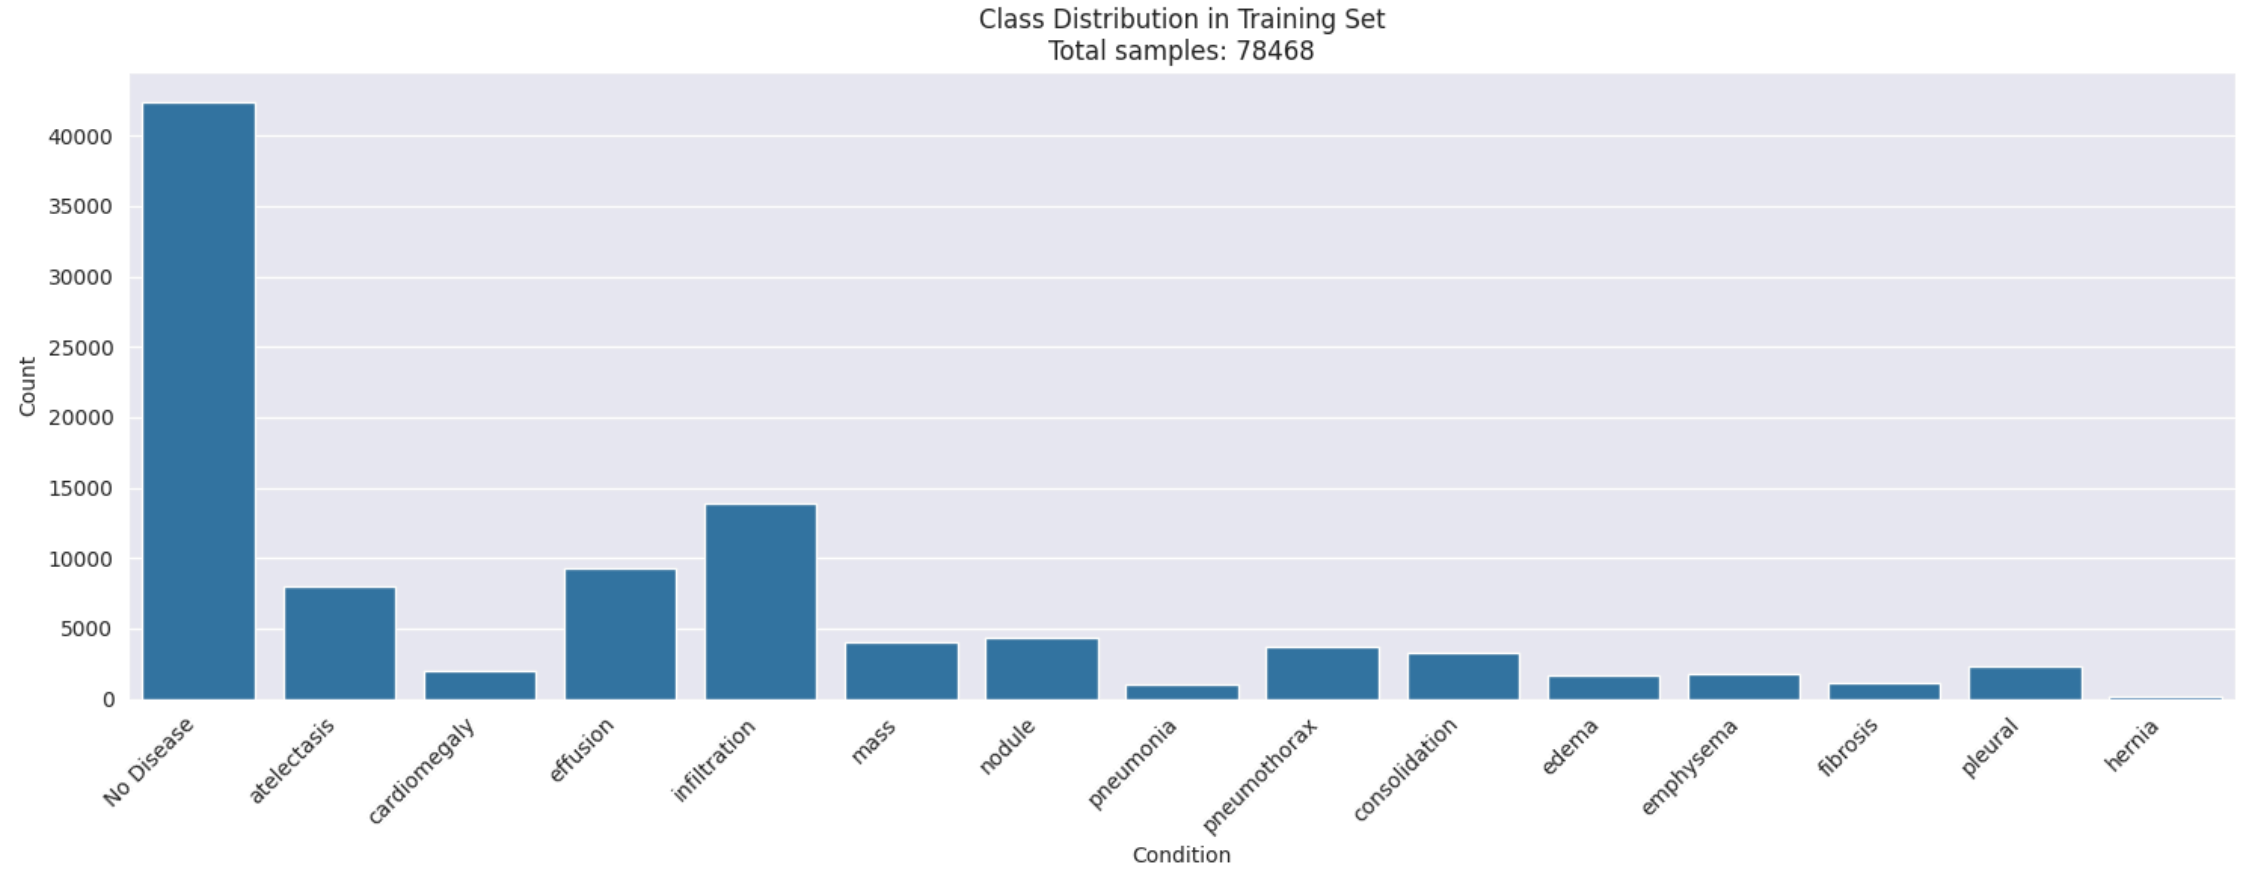
Medical Context

- Chest X-rays: Primary diagnostic tool
- Common conditions detected:
 - Pneumonia
 - Cardiomegaly
 - Edema
 - Pneumothorax
- Critical for rapid diagnosis and triage

Class distribution (unbalanced)



Class Distribution with Normal



Dataset Statistics

	Category	Count	Percentage
0	Total Samples	78468	
1	No Disease	42405	54.0%
2	Infiltration	13914	17.7%
3	Effusion	9261	11.8%
4	Atelectasis	7996	10.2%
5	Nodule	4375	5.6%
6	Mass	3988	5.1%
7	Pneumothorax	3705	4.7%
8	Consolidation	3263	4.2%
9	Pleural	2279	2.9%
10	Cardiomegaly	1950	2.5%
11	Emphysema	1799	2.3%
12	Edema	1690	2.2%
13	Fibrosis	1158	1.5%
14	Pneumonia	978	1.2%
15	Hernia	144	0.2%

Technical Architecture

Infrastructure

- PyTorch Lightning for training
- MLflow for experiment tracking
- Mixed precision training (FP16)
- GPU acceleration

Model Design

- Custom CNN architecture
- ResNet transfer learning option
- Early stopping and learning rate scheduling

Training Process

1. Data Pipeline

- Custom DataModule
- Augmentation strategies
- **×** Class weight balancing (wasn't sure about the best approach)

2. Training Loop

- Early stopping
- Learning rate scheduling
- Metrics monitoring
- Experiment tracking

Augmentation Strategies


- Preprocessing

```
transforms = A.Compose([
    A.RandomRotate90(p=0.5),
    A.VerticalFlip(p=0.5),
    A.ShiftScaleRotate(
        shift_limit=0.1,
        scale_limit=0.1,
        rotate_limit=rotate_limit,
        p=0.5,
    ),
    A.RandomBrightnessContrast(
        brightness_limit=brightness, contrast_limit=contrast, p=0.2
    ),
    A.Normalize(mean=[0.5], std=[0.5]),
    ToTensorV2(),
])
```

Condition:
Normal




Condition:
nodule




Condition:
Normal



Condition:
Normal




Condition:
Normal



Condition:
nodule



Condition:
Normal



Condition:
emphysema



Condition:
Normal



Condition:
edema



Custom CNN Architecture

```
class ChestNetS(ChestNetBase):
    """
    Simple CNN with three conv blocks and a classifier head. Each convolutional block includes convolution, batch normalization, ReLU activation,
    and max pooling operations. The classifier uses dropout for regularization and fully connected layers for final classification.
    Takes 64x64 grayscale images, outputs 14 binary classifications.
    """
    # Feature extraction backbone
    self.features = nn.Sequential(
        # Block 1: Input (1, 64, 64) -> Output (32, 32, 32)
        nn.Conv2d(1, 32, kernel_size=3, padding=1),
        nn.BatchNorm2d(32), # Normalize activations for stable training
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2, 2), # Reduce spatial dimensions by 2x


        # Block 2: Input (32, 32, 32) -> Output (64, 16, 16)
        nn.Conv2d(32, 64, kernel_size=3, padding=1),
        nn.BatchNorm2d(64),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2, 2),


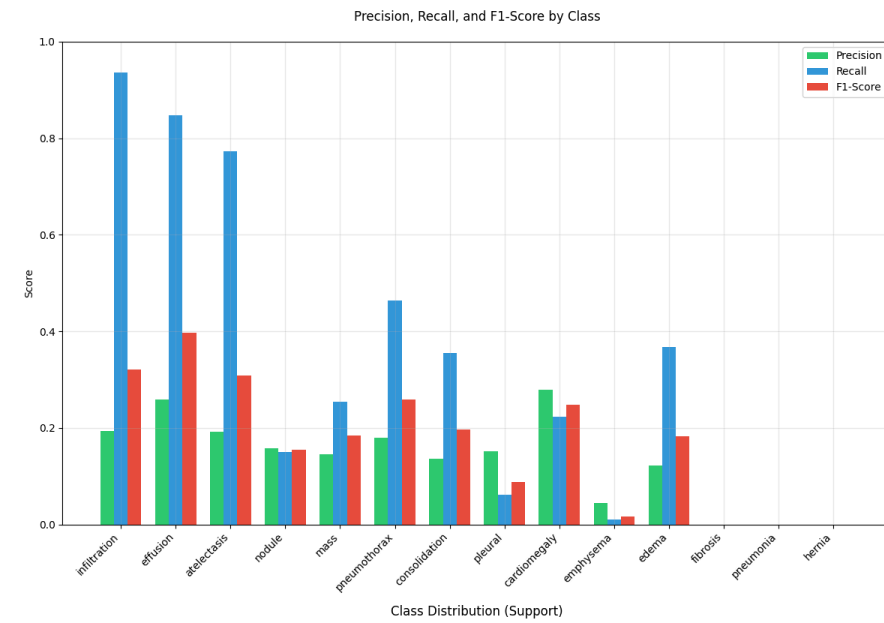
        # Block 3: Input (64, 16, 16) -> Output (128, 8, 8)
        nn.Conv2d(64, 128, kernel_size=3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2, 2),
    )

    # Classification head
    self.classifier = nn.Sequential(
        nn.Dropout(0.5), # Prevent overfitting
        nn.Linear(128 * 8 * 8, 512), # Flatten and project to 512 dimensions
        nn.ReLU(inplace=True),
        nn.Dropout(0.5), # Additional dropout layer
        nn.Linear(512, 14), # Final projection to output classes
    )


    def forward(self, x: torch.Tensor) -> torch.Tensor:
        x = self.features(x) # Extract visual features
        x = torch.flatten(x, 1) # Flatten spatial dimensions
        x = self.classifier(x) # Generate classification logits
        return x
```

CNN Metrics







Confusion Matrix - infiltration



Confusion Matrix - effusion







Original Image



GradCAM



Integrated Gradients






Resnet Architecture

```
class ChestNetResnet(ChestNetBase):
    # Input Size Compatibility: The ResNet-18 model is compatible with 64x64 inputs due to the
    # adaptive average pooling layer, which adjusts to varying spatial dimensions.
    # Pretrained Weights Handling: The first convolutional layer's weights are initialized by
    # averaging the pretrained RGB weights, preserving some pretrained features even with grayscale input.
    def __init__(
    ):
        # Load pretrained ResNet-18
        backbone = models.resnet18(pretrained=pretrained)
        # Modify first convolutional layer for grayscale input,
        original_conv1 = backbone.conv1
        self.backbone = backbone
        self.backbone.conv1 = nn.Conv2d(
            1, # Input channels changed to 1
            original_conv1.out_channels,
            kernel_size=original_conv1.kernel_size,
            stride=original_conv1.stride,
            padding=original_conv1.padding,
            bias=False,
        )
        # Initialize weights from pretrained model
        if pretrained:
            with torch.no_grad():
                self.backbone.conv1.weight.copy_(
                    original_conv1.weight.mean(dim=1, keepdim=True)
                )
        # Replace final fully connected layer
        in_features = self.backbone.fc.in_features
        self.backbone.fc = nn.Linear(in_features, num_classes)
        # Define feature extractor
        self.features = nn.Sequential(
            self.backbone.conv1,
            self.backbone.bn1,
            self.backbone.relu,
            self.backbone.maxpool,
            self.backbone.layer1,
            self.backbone.layer2,
            self.backbone.layer3,
            self.backbone.layer4,
            self.backbone.avgpool,
        )


    def forward(self, x: torch.Tensor) -> torch.Tensor:
        x = self.features(x)
        x = torch.flatten(x, 1)
        x = self.backbone.fc(x)
        return x
```

Resnet Metrics







Confusion Matrix - infiltration




Confusion Matrix - effusion







Original Image



GradCAM



Integrated Gradients



Hypothesis

- I think my early stopping might be flawed, Resnet early stopped at 40 epochs, while CNN at 125 epochs
- Class imbalance affecting the model performance
- I had to write statistics by hand, so I might have made a mistake
- 64x64 resolution might be too low for the model to learn the features (can a human learn from 64x64 images?)
- I didn't have time to implement hyperparameter optimization









Experiment Tracking



Experiments



ChestXRayPytorchLightning [Provide Feedback](#) [Add Description](#)

- ☐ Default  
- ☐ ChestXRayPytorchLightningRestNet18  
- ☒ ChestXRayPytorchLightning  
- ☐ ChestXRay  

Runs Evaluation **Experimental** Traces **Experimental**



Time created ▾

State: Active ▾

Datasets ▾

Sort: Created ▾

Columns ▾

 Group by ▾

<input type="checkbox"/>	Run Name	Created ▾	Dataset	Duration	Source	Models	Description
<input type="checkbox"/>	● Initial Run	✓ 2 days ago	-	1.3h	main.py	pytorch	Training run for chest X-ray classification using ...
<input type="checkbox"/>	● figuring out logging, che...	✗ 4 days ago	-	1.4d	main_n...	pytorch	Debug: Training run for chest X-ray classification...
<input type="checkbox"/>	● figuring out logging, che...	✓ 4 days ago	-	5.1min	main_n...	pytorch	Debug: Training run for chest X-ray classification...

ChestXRayPytorchLightningRestNet18 >

Initial Run


Register model

Overview **Model metrics** System metrics Artifacts

Search metric charts


Refresh

train_f1_score




Initial Run (train_f1_score)

val_precision




Initial Run (val_precision)

train_accuracy




Initial Run (train_accuracy)


wait_count



test_accuracy



val_f1_score



ChestXRayPytorchLightningRestNet18 >

Initial Run

Register model

Overview Model metrics System metrics **Artifacts**

- checkpoints
- model
 - data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - restored_model_checkpoint
 - auc_by_condition.png
 - class_metrics.png
 - condition_support.png
 - confusion_matrix_atelectasis.png
 - confusion_matrix_cardiomegaly.png
 - confusion_matrix_consolidation.png
 - confusion_matrix_edema.png
 - confusion_matrix_effusion.png
 - confusion_matrix_emphysema.png
 - confusion_matrix_fibrosis.png
 - confusion_matrix_hernia.png
 - confusion_matrix_infiltration.png
 - confusion_matrix_mass.png
 - confusion_matrix_nodule.png
 - confusion_matrix_pleural.png
 - confusion_matrix_pneumonia.png
 - confusion_matrix_pneumothorax.png
 - model_architecture.png
 - model_details.json
 - model_summary.txt
 - normal_detection.png
 - normal_detection_matrix.png
 - normal_detection_metrics.json
 - overall_metrics.json

model

Register model

Path: mlflow-artifacts:/738423306598801269/fd0c3cc9626a42ea83955cbaa349fc1f/artifacts/model

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. You can also [register it to the model registry](#) to version control

Model schema

Input and output schema for your model. [Learn more](#)

Name	Type
------	------

No schema. See [MLflow docs](#) for how to include input and output schema with your model.

Validate the model before deployment

Run the following code to validate model inference works on the example payload, prior to deploying it to a serving endpoint

```
from mlflow.models import validate_serving_input

model_uri = 'runs:/fd0c3cc9626a42ea83955cbaa349fc1f/model'

# The logged model does not contain an input_example.
# Manually generate a serving payload to verify your model prior to deployment.
from mlflow.models import convert_input_example_to_serving_input

# Define INPUT_EXAMPLE via assignment with your own input example to the model
# A valid input example is a data instance suitable for pyfunc prediction
serving_payload = convert_input_example_to_serving_input(INPUT_EXAMPLE)

# Validate the serving payload works on the model
validate_serving_input(model_uri, serving_payload)
```

Make Predictions

Predict on a Pandas DataFrame:

```
import mlflow
logged_model = 'runs:/fd0c3cc9626a42ea83955cbaa349fc1f/model'

# Load model as a PyFuncModel.
loaded_model = mlflow.pyfunc.load_model(logged_model)

# Predict on a Pandas DataFrame.
import pandas as pd
loaded_model.predict(pd.DataFrame(data))
```

ChestXRayPytorchLightningRestNet18 >

Initial Run


Register model

Overview Model metrics System metrics **Artifacts**

- checkpoints
- model
 - data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - restored_model_checkpoint
 - auc_by_condition.png**
 - class_metrics.png
 - condition_support.png
 - confusion_matrix_atelectasis.png
 - confusion_matrix_cardiomegaly.png
 - confusion_matrix_consolidation.png
 - confusion_matrix_edema.png
 - confusion_matrix_effusion.png
 - confusion_matrix_emphysema.png
 - confusion_matrix_fibrosis.png
 - confusion_matrix_hernia.png
 - confusion_matrix_infiltration.png
 - confusion_matrix_mass.png
 - confusion_matrix_nodule.png
 - confusion_matrix_pleural.png
 - confusion_matrix_pneumonia.png
 - confusion_matrix_pneumothorax.png
 - model_architecture.png
 - model_details.json
 - model_summary.txt
 - normal_detection.png
 - normal_detection_matrix.png
 - normal_detection_metrics.json
 - overall_metrics.json

auc_by_condition.png 36.51KB

Path: mlflow-artifacts:/738423306598801269/fd0c3cc9626a42ea83955cbaa349fc1f/artifacts/auc_by_condition.png



Intial Run


Register model

Overview Model metrics System metrics Artifacts

- checkpoints
- model
 - data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - restored_model_checkpoint
 - auc_by_condition.png
 - class_metrics.png
 - condition_support.png
 - confusion_matrix_atelectasis.png
 - confusion_matrix_cardiomegaly.png
 - confusion_matrix_consolidation.png
 - confusion_matrix_edema.png
 - confusion_matrix_effusion.png
 - confusion_matrix_emphysema.png
 - confusion_matrix_fibrosis.png
 - confusion_matrix_hernia.png
 - confusion_matrix_infiltration.png
 - confusion_matrix_mass.png
 - confusion_matrix_nodule.png
 - confusion_matrix_pleural.png
 - confusion_matrix_pneumonia.png
 - confusion_matrix_pneumothorax.png
 - model_architecture.png
 - model_details.json
 - model_summary.txt
 - normal_detection.png
 - normal_detection_matrix.png
 - normal_detection_metrics.json
 - overall_metrics.json

class_metrics.png 110.65KB

Path: mlflow-artifacts:/738423306598801269/fd0c3cc9626a42ea83955cbaa349fc1f/artifacts/class_metrics.png




Initial Run

Overview Model metrics System metrics **Artifacts**

- checkpoints
- model
 - data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - restored_model_checkpoint
 - auc_by_condition.png
 - class_metrics.png
 - condition_support.png
 - confusion_matrix_atelectasis.png**
 - confusion_matrix_cardiomegaly.png
 - confusion_matrix_consolidation.png
 - confusion_matrix_edema.png
 - confusion_matrix_effusion.png
 - confusion_matrix_emphysema.png
 - confusion_matrix_fibrosis.png
 - confusion_matrix_hernia.png
 - confusion_matrix_infiltration.png
 - confusion_matrix_mass.png
 - confusion_matrix_nodule.png
 - confusion_matrix_pleural.png
 - confusion_matrix_pneumonia.png
 - confusion_matrix_pneumothorax.png
 - model_architecture.png
 - model_details.json
 - model_summary.txt
 - normal_detection.png
 - normal_detection_matrix.png
 - normal_detection_metrics.json
 - overall_metrics.json

confusion_matrix_atelectasis.png 24.3KB

Path: milflow-artifacts:/738423306598801269/fd0c3cc9626a42ea83955cbaa349fc1f/artifacts/confusion_matrix_atelectasis.png



Initial Run

Overview Model metrics System metrics **Artifacts**

- ▶ checkpoints
- ▼ model
 - ▶ data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
 - restored_model_checkpoint
 - auc_by_condition.png
 - class_metrics.png
 - condition_support.png
 - confusion_matrix_atelectasis.png
 - confusion_matrix_cardiomegaly.png
 - confusion_matrix_consolidation.png
 - confusion_matrix_edema.png
 - confusion_matrix_effusion.png
 - confusion_matrix_emphysema.png
 - confusion_matrix_fibrosis.png
 - confusion_matrix_hernia.png
 - confusion_matrix_infiltration.png
 - confusion_matrix_mass.png
 - confusion_matrix_nodule.png
 - confusion_matrix_pleural.png
 - confusion_matrix_pneumonia.png
 - confusion_matrix_pneumothorax.png
 - model_architecture.png
 - model_details.json
 - model_summary.txt**
 - normal_detection.png
 - normal_detection_matrix.png
 - normal_detection_metrics.json
 - overall_metrics.json

model_summary.txt 6.35KBPath: milflow-artifacts:/738423306598801269/fd0c3cc9626a42ea83955cbaa349fc1f/artifacts/model_summary.txt [📄](#)

	Name	Type	Params	Mode
0	train_metrics	MetricCollection	0	train
1	train_metrics.accuracy	MultilabelAccuracy	0	train
2	train_metrics.f1_score	MultilabelF1Score	0	train
3	train_metrics.precision	MultilabelPrecision	0	train
4	val_metrics	MetricCollection	0	train
5	val_metrics.accuracy	MultilabelAccuracy	0	train
6	val_metrics.f1_score	MultilabelF1Score	0	train
7	val_metrics.precision	MultilabelPrecision	0	train
8	test_metrics	MetricCollection	0	train
9	test_metrics.accuracy	MultilabelAccuracy	0	train
10	test_metrics.f1_score	MultilabelF1Score	0	train
11	test_metrics.precision	MultilabelPrecision	0	train
12	backbone	ResNet	11.2 M	train
13	backbone.conv1	Conv2d	3.1 K	train
14	backbone.bn1	BatchNorm2d	128	train
15	backbone.relu	ReLU	0	train
16	backbone.maxpool	MaxPool2d	0	train
17	backbone.layer1	Sequential	147 K	train
18	backbone.layer1.0	BasicBlock	74.0 K	train
19	backbone.layer1.0.conv1	Conv2d	36.9 K	train
20	backbone.layer1.0.bn1	BatchNorm2d	128	train
21	backbone.layer1.0.relu	ReLU	0	train
22	backbone.layer1.0.conv2	Conv2d	36.9 K	train
23	backbone.layer1.0.bn2	BatchNorm2d	128	train
24	backbone.layer1.1	BasicBlock	74.0 K	train
25	backbone.layer1.1.conv1	Conv2d	36.9 K	train
26	backbone.layer1.1.bn1	BatchNorm2d	128	train
27	backbone.layer1.1.relu	ReLU	0	train
28	backbone.layer1.1.conv2	Conv2d	36.9 K	train
29	backbone.layer1.1.bn2	BatchNorm2d	128	train
30	backbone.layer2	Sequential	525 K	train
31	backbone.layer2.0	BasicBlock	230 K	train
32	backbone.layer2.0.conv1	Conv2d	73.7 K	train
33	backbone.layer2.0.bn1	BatchNorm2d	256	train
34	backbone.layer2.0.relu	ReLU	0	train
35	backbone.layer2.0.conv2	Conv2d	147 K	train

Intial Run

Register model

Overview Model metrics System metrics Artifacts

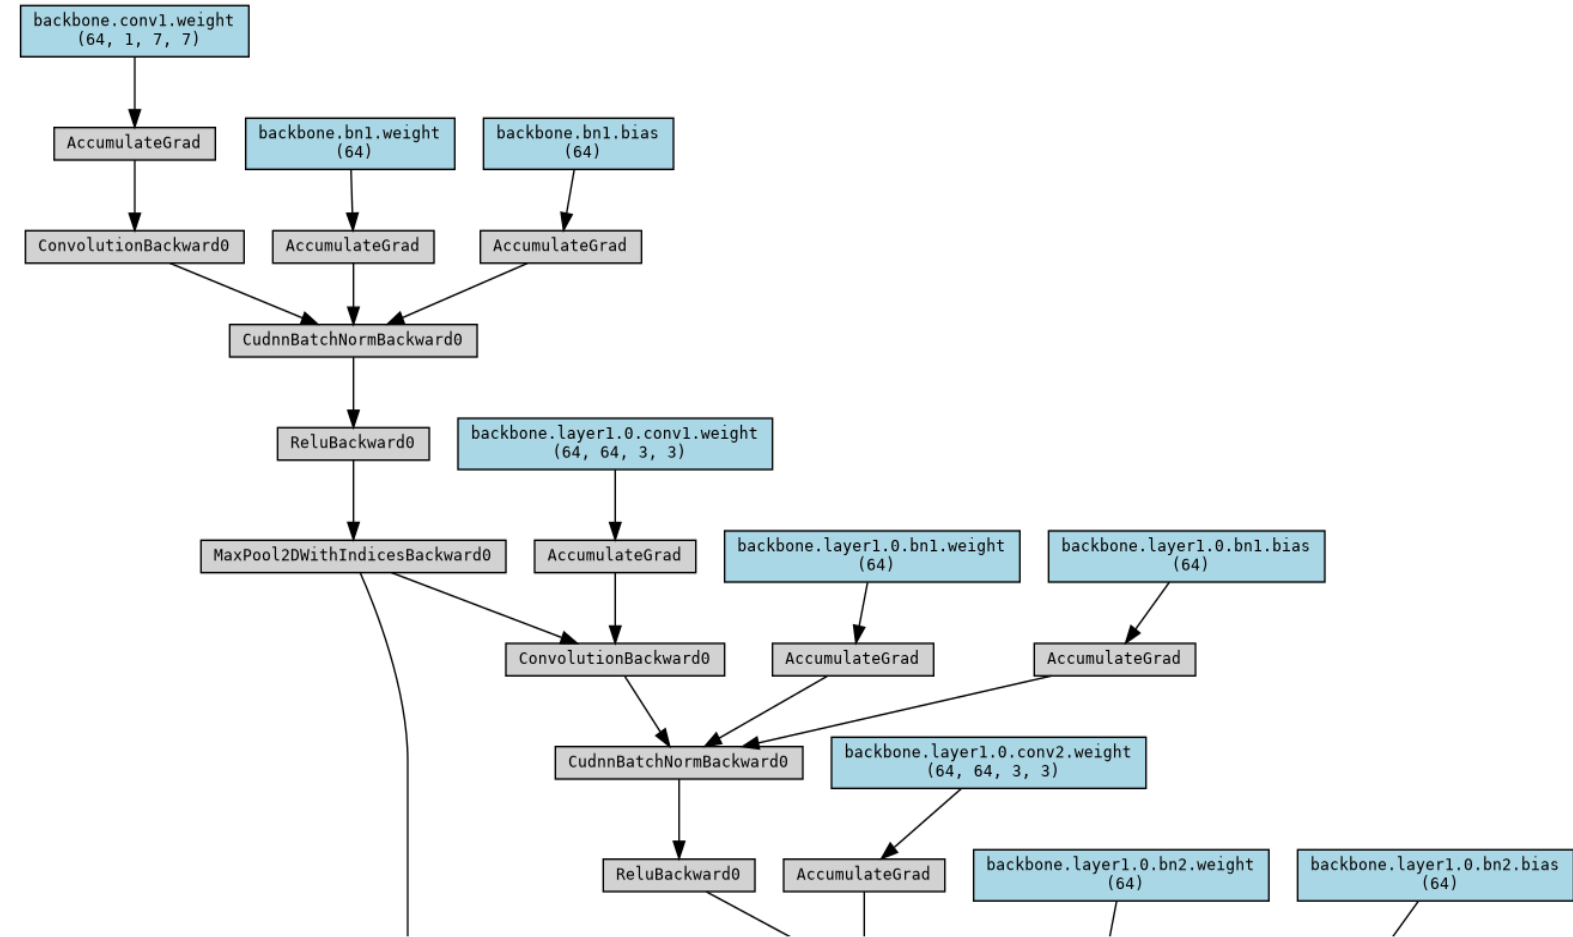
checkpoints

model

- data
 - MLmodel
 - conda.yaml
 - python_env.yaml
 - requirements.txt
- restored_model_checkpoint
 - auc_by_condition.png
 - class_metrics.png
 - condition_support.png
 - confusion_matrix_atelectasis.png
 - confusion_matrix_cardiomegaly.png
 - confusion_matrix_consolidation.png
 - confusion_matrix_edema.png
 - confusion_matrix_effusion.png
 - confusion_matrix_emphysema.png
 - confusion_matrix_fibrosis.png
 - confusion_matrix_hernia.png
 - confusion_matrix_infiltration.png
 - confusion_matrix_mass.png
 - confusion_matrix_nodule.png
 - confusion_matrix_pleural.png
 - confusion_matrix_pneumonia.png
 - confusion_matrix_pneumothorax.png
 - model_architecture.png
 - model_details.json
 - model_summary.txt
 - normal_detection.png
 - normal_detection_matrix.png
 - normal_detection_metrics.json
 - overall_metrics.json

model_architecture.png 862.21KB

Path: mlflow-artifacts:/738423306598801269/fd0c3cc9626a42ea83955cbaa349fc1f/artifacts/model_architecture.png



Thank You

Questions?

Contact: piotr.gryko@gmail.com